

---

---









---

# NoSQL Databases

Principles of Database Management

Lemahieu et. Al.

Cambridge University Press, 2018

---

# Introduction

---

- The NoSQL movement
- Key–value stores
- Tuple and document stores
- Column-oriented databases
- Graph-based databases
- Other NoSQL categories

# The NoSQL Movement

---

- RDBMSs put a lot of emphasis on keeping data consistent.
  - Entire database is consistent at all times (ACID)
- Focus on consistency may hamper flexibility and scalability
- As the data volumes or number of parallel transactions increase, capacity can be increased by
  - Vertical scaling: extending storage capacity and/or CPU power of the database server
  - Horizontal scaling: multiple DBMS servers being arranged in a cluster



# The NoSQL Movement

---

- RDBMSs are not good at extensive horizontal scaling
  - Coordination overhead because of focus on consistency
  - Rigid database schemas
- Other types of DBMSs needed for situations with massive volumes, flexible data structures, and where scalability and availability are more important → NoSQL databases

# The NoSQL Movement

---

- NoSQL databases
  - Describes databases that store and manipulate data in formats other than tabular relations, i.e., non-relational databases (NoREL)
- NoSQL databases aim at near-linear horizontal scalability by distributing data over a cluster of database nodes for the sake of performance as well as availability
- Eventual consistency: the data (and its replicas) will become consistent at some point in time after each transaction

# The NoSQL Movement

	Relational Databases	NoSQL Databases
<b>Data paradigm</b>	Relational tables	Key–value (tuple) based Document based Column based Graph based XML, object based Others: time series, probabilistic, etc.
<b>Distribution</b>	Single-node and distributed	Mainly distributed
<b>Scalability</b>	Vertical scaling, harder to scale horizontally	Easy to scale horizontally, easy data replication
<b>Openness</b>	Closed and open source	Mainly open source
<b>Schema role</b>	Schema-driven	Mainly schema-free or flexible schema
<b>Query language</b>	SQL as query language	No or simple querying facilities, or special-purpose languages
<b>Transaction mechanism</b>	ACID: Atomicity, Consistency, Isolation, Durability	BASE: Basically Available, Soft state, Eventual consistency
<b>Feature set</b>	Many features (triggers, views, stored procedures, etc.)	Simple API
<b>Data volume</b>	Capable of handling normal-sized datasets	Capable of handling huge amounts of data and/or very high frequencies of read/write requests

# Key–Value Stores

---

- Key–value-based database stores data as (key, value) pairs
  - Keys are unique
  - Hash map, or hash table or dictionary

# Key–Value Stores

---

```
import java.util.HashMap;
import java.util.Map;
public class KeyValueStoreExample {
    public static void main(String... args) {
        // Keep track of age based on name
        Map<String, Integer> age_by_name = new HashMap<>();

        // Store some entries
        age_by_name.put("wilfried", 34);
        age_by_name.put("seppe", 30);
        age_by_name.put("bart", 46);
        age_by_name.put("jeanne", 19);

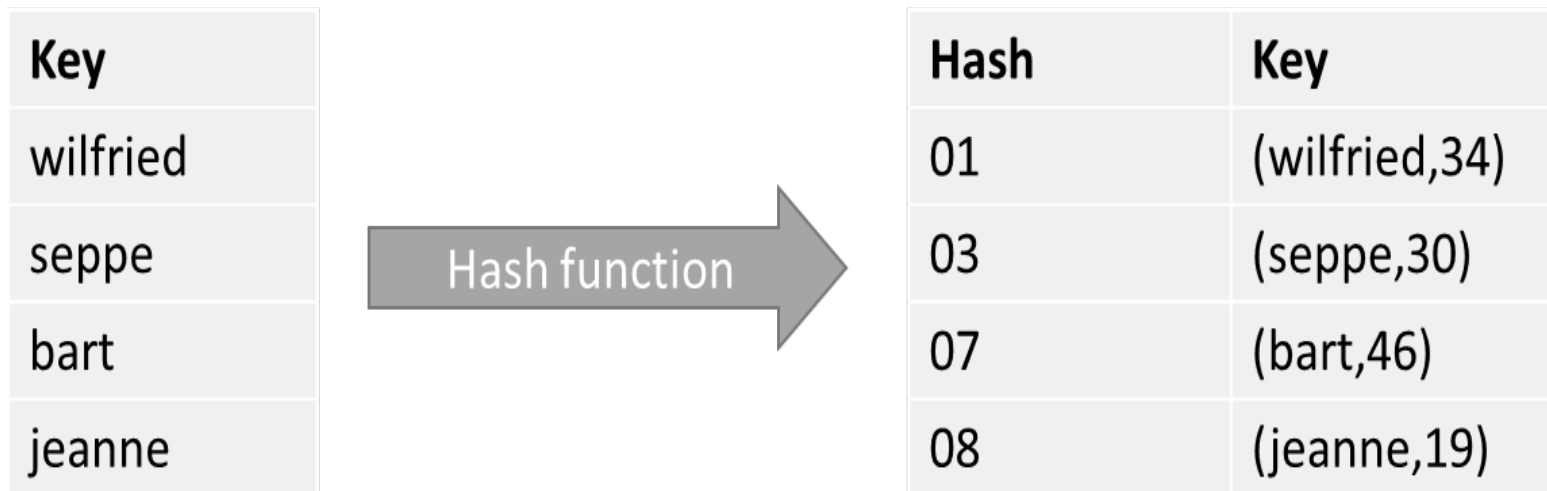
        // Get an entry
        int age_of_wilfried = age_by_name.get("wilfried");
        System.out.println("Wilfried's age: " + age_of_wilfried);

        // Keys are unique
        age_by_name.put("seppe", 50); // Overrides previous entry
    }
}
```

# Key–Value Stores

---

- Keys (e.g., “bart”, “seppe”) are hashed by means of a so-called **hash function**
  - A hash function takes an arbitrary value of arbitrary size and maps it to a key with a fixed size, which is called the hash value
  - Each hash can be mapped to a space in computer memory

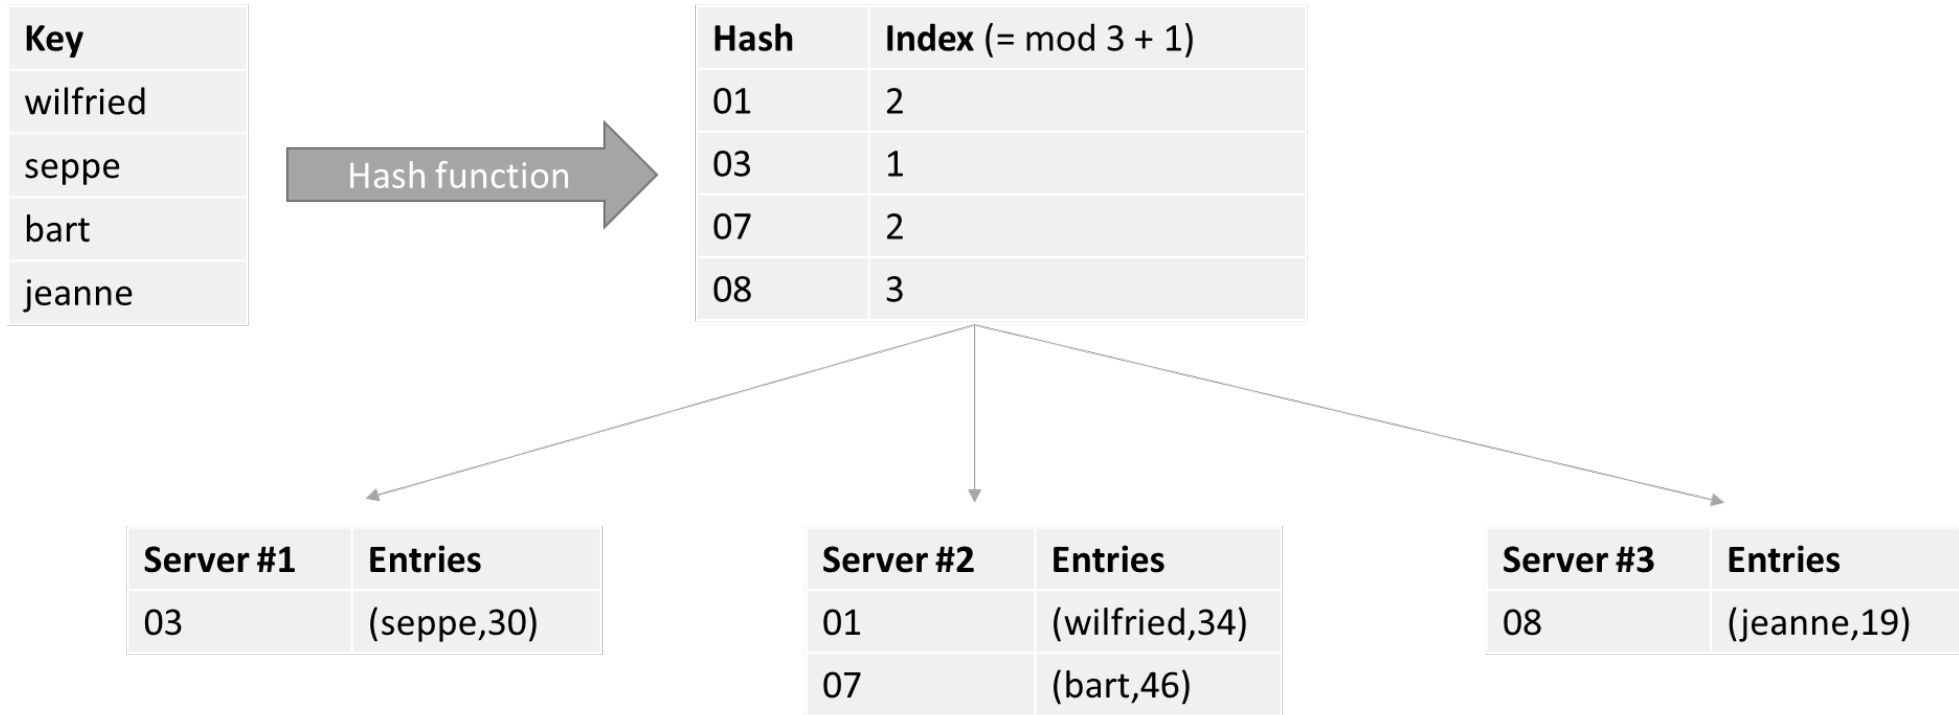


# Key–Value Stores

---

- NoSQL databases are built with horizontal scalability support in mind
- Distribute hash table over different locations
- Assume we need to spread our hashes over three servers
  - Hash every key (“wilfried”, “seppe”) to a server identifier
  - $\text{index}(\text{hash}) = \text{mod}(\text{hash}, \text{nrServers}) + 1$

# Key-Value Stores



Sharding!



# Key–Value Stores

---

- Example: Memcached
  - Implements a distributed memory-driven hash table (i.e., a key–value store), which is put in front of a traditional database to speed up queries by caching recently accessed objects in RAM
  - Caching solution

# Key–Value Stores

---

```
import java.util.ArrayList;
import java.util.List;
import net.spy.memcached.AddrUtil;
import net.spy.memcached.MemcachedClient;

public class MemCachedExample {
    public static void main(String[] args) throws Exception {
        List<String> serverList = new ArrayList<String>() {
            {
                this.add("memcachedserver1.servers:11211");
                this.add("memcachedserver2.servers:11211");
                this.add("memcachedserver3.servers:11211");
            }
        };
    }
};
```

# Key–Value Stores

---

```
MemcachedClient memcachedClient = new MemcachedClient(
    AddrUtil.getAddresses(serverList));

// ADD adds an entry and does nothing if the key already exists
// Think of it as an INSERT
// The second parameter (0) indicates the expiration - 0 means no expiry
memcachedClient.add("marc", 0, 34);
memcachedClient.add("seppe", 0, 32);
memcachedClient.add("bart", 0, 66);
memcachedClient.add("jeanne", 0, 19);

// SET sets an entry regardless of whether it exists
// Think of it as an UPDATE-OR-INSERT
memcachedClient.add("marc", 0, 1111); // <- ADD will have no effect
memcachedClient.set("jeanne", 0, 12); // <- But SET will
```

# Key–Value Stores

---

```
// REPLACE replaces an entry and does nothing if the key does not exist
// Think of it as an UPDATE
memcachedClient.replace("not_existing_name", 0, 12); // <- Will have no effect
memcachedClient.replace("jeanne", 0, 10);

// DELETE deletes an entry, similar to an SQL DELETE statement
memcachedClient.delete("seppe");

// GET retrieves an entry
Integer age_of_marc = (Integer) memcachedClient.get("marc");
Integer age_of_short_lived = (Integer) memcachedClient.get("short_lived_name");
Integer age_of_not_existing = (Integer) memcachedClient.get("not_existing_name");
Integer age_of_seppe = (Integer) memcachedClient.get("seppe");
System.out.println("Age of Marc: " + age_of_marc);
System.out.println("Age of Seppe (deleted): " + age_of_seppe);
System.out.println("Age of not existing name: " + age_of_not_existing);
System.out.println("Age of short lived name (expired): " + age_of_short_lived);

memcachedClient.shutdown();

}
}
```

# Key–Value Stores

---

- Request coordination
- Consistent hashing
- Replication and redundancy
- Eventual consistency
- Stabilization
- Integrity constraints and querying

# Request Coordination

---

- In many NoSQL implementations (e.g., Cassandra, Google's BigTable, Amazon's DynamoDB), all nodes implement the same functionality and are all able to perform the role of request coordinator
- Need for membership protocol
  - Dissemination
    - Based on periodic, pairwise communication
  - Failure detection

# Consistent Hashing

---

- **Consistent hashing** schemes are often used, which avoid having to remap each key to a new node when nodes are added or removed
- Suppose we have a situation in which ten keys are distributed over three servers ( $n = 3$ ) with the following hash function:
  - $h(key) = key \text{ modulo } n$

# Consistent Hashing

---


	<i>n</i>		
key	3	2	4
0	0	0	0
1	1	1	1
2	2	0	2
3	0	1	3
4	1	0	0
5	2	1	1
6	0	0	2
7	1	1	3
8	2	0	0
9	0	1	1



# Consistent Hashing


---

- At the core of a consistent hashing setup is a so-called **“ring”-topology**, which is basically a representation of the number range  $[0,1]$ :




# Consistent Hashing

---



# Consistent Hashing

- Hash each key to a position on the ring, and store the actual key–value pair on the first server that appears clockwise of the hashed point on the ring




# Consistent Hashing

---

- Because of the uniformity property of a “good” hash function, roughly  $1/n$  of key–value pairs will end up being stored on each server
- Most of the key–value pairs will remain unaffected in the event that a machine is added or removed

# Consistent Hashing



# Replication and Redundancy

---

- Problems with consistent hashing:
  - If two servers end up being mapped close to one another, one of these nodes will end up with few keys to store
  - In the case that a server is added, all of the keys moved to this new node originate from just one other server
- Instead of mapping a server  $s$  to a single point on our ring, we map it to multiple positions, called **replicas**
- For each physical server  $s$ , we hence end up with  $r$  (the number of replicas) points on the ring
- Note: each of the replicas still represents the same physical instance ( $\leftrightarrow$  redundancy)
  - Virtual nodes


# Replication and Redundancy

---

- To handle data replication or redundancy, many vendors extend the consistent hashing mechanism so that key–value pairs are duplicated across multiple nodes
  - e.g., by storing the key–value pair on two or more nodes clockwise from the key's position on the ring

# Replication and Redundancy

- It is also possible to set up a full redundancy scheme in which each node itself corresponds to multiple physical machines, each storing a fully redundant copy of the data





# Eventual Consistency

---

- Membership protocol does not guarantee that every node is aware of every other node *at all times*
  - it will reach a consistent state over time
- State of the network might not be perfectly consistent at any moment in time, though will become eventually consistent at a future point in time
- Many NoSQL databases guarantee so-called **eventual consistency**

# Eventual Consistency

---

- Most NoSQL databases follow the **BASE** principle
  - Basically Available, Soft state, Eventual consistency
- **CAP theorem** states that a distributed computer system cannot guarantee the following three properties at the same time:
  - Consistency (all nodes see the same data at the same time)
  - Availability (guarantees that every request receives a response indicating a success or failure result)
  - Partition tolerance (the system continues to work even if nodes go down or are added)

# Eventual Consistency

---

- Most NoSQL databases sacrifice the consistency part of CAP in their setup, instead striving for eventual consistency
- The full BASE acronym stands for:
  - Basically Available: NoSQL databases adhere to the availability guarantee of the CAP theorem
  - Soft state: the system can change over time, even without receiving input
  - Eventual consistency: the system will become consistent over time

# Stabilization

---

- The operation which repartitions hashes over nodes in case nodes are added or removed is called **stabilization**
- If a consistent hashing scheme is being applied, the number of fluctuations in the hash–node mappings will be minimized.

# Integrity Constraints and Querying

---

- Key–value stores represent a very diverse gamut of systems
- Full-blown DBMSs versus caches
- Only limited query facilities are offered
  - e.g. put and set
- Limited to no means to enforce structural constraints
  - DBMS remains agnostic to the internal structure
- No relationships, referential integrity constraints, or database schema can be defined

# Tuple and Document Stores

---

- A **tuple store** is similar to a key–value store, with the difference that it does not store pairwise combinations of a key and a value, but instead stores a unique key together with a vector of data
- Example:
  - marc -> ("Marc", "McLast Name", 25, "Germany")
- No requirement to have the same length or semantic ordering (schema-less!)

# Tuple and Document Stores

---

- Various NoSQL implementations do, however, permit organizing entries in semantical groups (aka collections or tables)
- Examples:
  - `Person:marc -> ("Marc", "McLast Name", 25, "Germany")`
  - `Person:harry -> ("Harry", "Smith", 29, "Belgium")`

# Tuple and Document Stores

---

- **Document stores** store a collection of attributes that are labeled and unordered, representing items that are semi-structured

- Example:

```
{  
  Title = "Harry Potter"  
  ISBN = "111-1111111111"  
  Authors = [ "J.K. Rowling" ]  
  Price = 32  
  Dimensions = "8.5 x 11.0 x 0.5"  
  PageCount = 234  
  Genre = "Fantasy"  
}
```



# Tuple and Document Stores

---

- Most modern NoSQL databases choose to represent documents using JSON

```
{  
  "title": "Harry Potter",  
  "authors": ["J.K. Rowling", "R.J. Kowling"],  
  "price": 32.00,  
  "genres": ["fantasy"],  
  "dimensions": {  
    "width": 8.5,  
    "height": 11.0,  
    "depth": 0.5  
  },  
  "pages": 234,  
  "in_publication": true,  
  "subtitle": null  
}
```

# Tuple and Document Stores

---

- Items with keys
- Filters and queries
- Complex queries and aggregation with MapReduce
- SQL after all ...

# Items with Keys

---

- Most NoSQL document stores will allow you to store items in tables (collections) in a schema-less manner, but will enforce that a primary key be specified
  - e.g. Amazon's DynamoDB, MongoDB ( `_id` )
- A primary key will be used as a partitioning key to create a hash and determine where the data will be stored

# Filters and Queries

---

```
import org.bson.Document;
import com.mongodb.MongoClient;
import com.mongodb.client.FindIterable;
import com.mongodb.client.MongoDatabase;
import java.util.ArrayList;
import static com.mongodb.client.model.Filters.*;
import static java.util.Arrays.asList;

public class MongoDBExample {
    public static void main(String... args) {
        MongoClient mongoClient = new MongoClient();
        MongoDatabase db = mongoClient.getDatabase("test");

        // Delete all books first
        db.getCollection("books").deleteMany(new Document());

        // Add some books
        db.getCollection("books").insertMany(new ArrayList<Document>() {{
            add(getBookDocument("My First Book", "Wilfried", "Lemahieu", 12, new String[]{"drama"}));
            add(getBookDocument("My Second Book", "Seppe", "vanden Broucke", 437, new String[]{"fantasy", "thriller"}));
            add(getBookDocument("My Third Book", "Seppe", "vanden Broucke", 200, new String[]{"educational"}));
            add(getBookDocument("Java Programming", "Bart", "Baesens", 100, new String[]{"educational"}));
        }});
```

# Filters and Queries

---

```
// Perform query
FindIterable<Document> result = db.getCollection("books").find(
    and( eq("author.last_name", "vanden Broucke"),
        eq("genres", "thriller"),
        gt("nrPages", 100)));

for (Document r : result) {
    System.out.println(r.toString());
    // Increase the number of pages:
    db.getCollection("books").updateOne(
        new Document("_id", r.get("_id")),
        new Document("$set",
            new Document("nrPages", r.getInteger("nrPages") + 100)));
}
mongoClient.close();}
```

```
public static Document getBookDocument(String title,
    String authorFirst, String authorLast,
    int nrPages, String[] genres) {
    return new Document("author", new Document()
        .append("first_name", authorFirst)
        .append("last_name", authorLast))
        .append("title", title)
        .append("nrPages", nrPages)
        .append("genres", asList(genres)));}
```

# Filters and Queries

---

Document({\_id=567ef62bc0c3081f4c04b16c,  
author=Document{{first\_name=Seppe, last\_name=vanden Broucke}},  
title=My Second Book, nrPages=437, genres=[fantasy, thriller]}}

# Filters and Queries

---

```
// Perform aggregation query
AggregateIterable<Document> result = db.getCollection("books")
    .aggregate(asList(
        new Document("$group",
            new Document("_id", "$author.last_name")
                .append("page_sum", new Document("$sum",
                    "$nrPages")))));

for (Document r : result) {
    System.out.println(r.toString());
}
```

```
Document{{_id=Lemahieu, page_sum=12}}
Document{{_id=Vanden Broucke, page_sum=637}}
Document{{_id=Baesens, page_sum=100}}
```

# Filters and Queries

---

- Queries can still be slow because every filter (such as “author.last\_name = Baesens”) entails a complete collection or table scan
- Most document stores can define a variety of indexes
  - unique and non-unique indexes
  - compound indexes
  - geospatial indexes
  - text-based indexes



# Complex Queries and Aggregation with MapReduce

---

- Document stores do not support relations
- **First approach:** embedded documents

```
{  
  "title": "Databases for Beginners",  
  "authors": ["J.K. Sequel", "John Smith"],  
  "pages": 234  
}
```

**BUT: Data duplication!**

```
{  
  "title": "Databases for Beginners",  
  "authors": [  
    {"first_name": "Jay Kay", "last_name": "Sequel", "age": 54},  
    {"first_name": "John", "last_name": "Smith", "age": 32}  
  ],  
  "pages": 234  
}
```

# Complex Queries and Aggregation with MapReduce

---

- **Second approach:** create two collections

book collection:

```
{  
  "title": "Databases for Beginners",  
  "authors": ["Jay Kay Rowling", "John Smith"],  
  "pages": 234  
}
```

authors collection:

```
{  
  "_id": "Jay Kay Rowling",  
  "age": 54  
}
```

**BUT: Need to resolve complex relational queries in application code!**

# Complex Queries and Aggregation with MapReduce

---

- **Third Approach: MapReduce**

- A map-reduce pipeline starts from a series of key–value pairs  $(k_1, v_1)$  and maps each pair to one or more output pairs
- The output entries are shuffled and distributed so that all output entries belonging to the same key are assigned to the same worker (e.g., physical machines)
- Workers then apply a reduce function to each group of key–value pairs having the same key, producing a new list of values per output key
- The resulting, final outputs are then (optionally) sorted per key  $k_2$  to produce the final outcome

# Complex Queries and Aggregation with MapReduce

---

- Example: get a summed count of pages for books per genre
- Create a list of input key–value pairs

k1	v1
1	{genre: education, nrPages: 120}
2	{genre: thriller, nrPages: 100}
3	{genre: fantasy, nrPages: 20}
...	...

- Map function is a simple conversion to a genre–nrPages key–value pair

```
function map(k1, v1)  
    emit output record (v1.genre, v1.nrPages)  
end function
```

# Complex Queries and Aggregation with MapReduce

---

- Workers have produced the following three output lists, with the keys corresponding to genres

Worker 1	
k2	v2
education	120
thriller	100
fantasy	20

Worker 2	
k2	v2
drama	500
education	200

Worker 3	
k2	v2
education	20
fantasy	10

- A working operation will be started per unique key k2, for which its associated list of values will be reduced
  - e.g., (education,[120,200,20]) will be reduced to its sum, 340

```
function reduce(k2, v2_list)
    emit output record (k2, sum(v2_list))
end function
```

# Complex Queries and Aggregation with MapReduce

---

- Final output looks like this:

k2	v3
education	340
thriller	100
drama	500
fantasy	30

- Can be sorted based on k2 or v3

# Complex Queries and Aggregation with MapReduce

---

- Suppose we would now like to retrieve an average page count per book for each genre
- Reduce function becomes

```
function reduce(k2, v2_list)
    emit output record (k2, sum(v2_list) / length(v2_list))
end function
```
- After mapping the input list, workers produce the following three output lists

Worker 1	
k2	v2
education	120
thriller	100
fantasy	20

Worker 2	
k2	v2
drama	500
education	200

Worker 3	
k2	v2
education	20
fantasy	10

# Complex Queries and Aggregation with MapReduce

---

- Average as follows

k2	v3
education	$(20 + 50 + 50 + 100 + 100 + 20) / 6 = 56.67$
thriller	$100 / 1 = 100.00$
drama	$(100 + 200 + 200) / 3 = 166.67$
fantasy	$(20 + 10) / 3 = 10.00$

- Note: reduce operation can happen more than once, and can already start before all mapping operations have finished!
  - Need to ensure that results are correct by possibly rewriting map and reduce functions!



# Complex Queries and Aggregation with MapReduce

---

- Example: count the number of occurrences per word in a document

```
function map(document_name, document_text)
    for each word in document_text do
        emit output record (word, 1)
    repeat
end function
```

```
function reduce(word, partial_counts)
    emit output record (word, sum(partial_counts))
end function
```

# Complex Queries and Aggregation with MapReduce

---

- Example: return the average number of pages per genre, but now taking into account that books can have more than one genre associated with them (in MongoDB)

# Complex Queries and Aggregation with MapReduce

---

```
import org.bson.Document;
import com.mongodb.MongoClient;
import com.mongodb.client.MongoDatabase;
import java.util.ArrayList;
import java.util.List;
import java.util.Random;

import static java.util.Arrays.asList;

public class MongoDBAggregationExample {
    public static Random r = new Random();

    public static void main(String... args) {
        MongoClient mongoClient = new MongoClient();
        MongoDatabase db = mongoClient.getDatabase("test");

        setupDatabase(db);
        for (Document r : db.getCollection("books").find())
            System.out.println(r);

        mongoClient.close();}
```

# Complex Queries and Aggregation with MapReduce

---

```
public static void setupDatabase(MongoDatabase db) {
    db.getCollection("books").deleteMany(new Document());

    String[] possibleGenres = new String[] {
        "drama", "thriller", "romance", "detective",
        "action", "educational", "humor", "fantasy" };

    for (int i = 0; i < 100; i++) {
        db.getCollection("books").insertOne(
            new Document("_id", i)
            .append("nrPages", r.nextInt(900) + 100)
            .append("genres",
getRandom(asList(possibleGenres), r.nextInt(3) + 1)));
    }
}
```

# Complex Queries and Aggregation with MapReduce

---

```
public static List<String> getRandom(List<String> els, int number) {  
    List<String> selected = new ArrayList<>();  
    List<String> remaining = new ArrayList<>(els);  
    for (int i = 0; i < number; i++) {  
        int s = r.nextInt(remaining.size());  
        selected.add(remaining.get(s));  
        remaining.remove(s);  
    }  
    return selected;  
}  
}
```

Document{{\_id=0, nrPages=188, genres=[action, detective, romance]}}

Document{{\_id=1, nrPages=976, genres=[romance, detective, humor]}}

Document{{\_id=2, nrPages=652, genres=[thriller, fantasy, action]}}

Document{{\_id=3, nrPages=590, genres=[fantasy]}}

Document{{\_id=4, nrPages=703, genres=[educational, drama, thriller]}}

Document{{\_id=5, nrPages=913, genres=[detective]}}

...

# Complex Queries and Aggregation with MapReduce

- Manual construction of the aggregation query looks as follows:

```
public static void reportAggregate(MongoDatabase db) {
    Map<String, List<Integer>> counts = new HashMap<>();
    for (Document r : db.getCollection("books").find()) {
        for (Object genre : r.get("genres", List.class)) {
            if (!counts.containsKey(genre.toString()))
                counts.put(genre.toString(), new ArrayList<Integer>());
            counts.get(genre.toString()).add(r.getInteger("nrPages"));
        }
    }
    for (Entry<String, List<Integer>> entry : counts.entrySet()) {
        System.out.println(entry.getKey() + " --> AVG = " +
            sum(entry.getValue()) / (double) entry.getValue().size());
    }
}

private static int sum(List<Integer> value) {
    int sum = 0;
    for (int i : value) sum += i;
    return sum;
}
```

romance --> AVG = 497.39285714285717  
drama --> AVG = 536.88  
detective --> AVG = 597.1724137931035  
humor --> AVG = 603.5357142857143  
fantasy --> AVG = 540.0434782608696  
educational --> AVG = 536.1739130434783  
action --> AVG = 398.9032258064516  
thriller --> AVG = 513.5862068965517

# Complex Queries and Aggregation with MapReduce

---

- If the list of genres is known beforehand, we can optimize by performing the aggregation per genre directly in MongoDB itself:

```
public static void reportAggregate(MongoDatabase db) {
    String[] possibleGenres = new String[] {
        "drama", "thriller", "romance", "detective",
        "action", "educational", "humor", "fantasy" };

    for (String genre : possibleGenres) {
        AggregateIterable<Document> iterable =
            db.getCollection("books").aggregate(asList(
                new Document("$match", new Document("genres", genre)),
                new Document("$group", new Document("_id", genre)
                    .append("average", new Document("$avg", "$nrPages"))));

        for (Document r : iterable) {
            System.out.println(r);
        }
    }
}
```

Document{[\_id=drama, average=536.88]}  
Document{[\_id=thriller, average=513.5862068965517]}  
Document{[\_id=romance, average=497.39285714285717]}  
Document{[\_id=detective, average=597.1724137931035]}  
Document{[\_id=action, average=398.9032258064516]}  
...

# Complex Queries and Aggregation with MapReduce

---

- Assume now that we have millions of books in our database and we do not know the number of genres beforehand → use MapReduce
- Map in MongoDB:

```
function() {  
    // No arguments, use “this” to refer to the  
    // local document item being processed  
    emit(key, value);  
}
```

- Reduce in MongoDB:

```
function(key, values) {  
    return result;  
}
```



# Complex Queries and Aggregation with MapReduce

---

- **Map function**

```
function() {  
    var nrPages = this.nrPages;  
    this.genres.forEach(function(genre) {  
        emit(genre, {average: nrPages, count: 1});  
    });  
}
```

- **Reduce function**

```
function(genre, values) {  
    var s = 0;  
    var newc = 0;  
    values.forEach(function(curAvg) {  
        s += curAvg.average * curAvg.count;  
        newc += curAvg.count;  
    });  
    return {average: (s / newc), count: newc};  
}
```

# Complex Queries and Aggregation with MapReduce

---

```
public static void reportAggregate(MongoDatabase db) {
    String map = "function() { " +
        " var nrPages = this.nrPages; " +
        " this.genres.forEach(function(genre) { " +
        " emit(genre, {average: nrPages, count: 1}); " +
        " }); " +
        "} ";

    String reduce = "function(genre, values) { " +
        " var s = 0; var newc = 0; " +
        " values.forEach(function(curAvg) { " +
        " s += curAvg.average * curAvg.count; " +
        " newc += curAvg.count; " +
        " }); " +
        " return {average: (s / newc), count: newc}; " +
        "} ";

    MapReduceIterable<Document> result = db.getCollection("books")
        .mapReduce(map, reduce);
    for (Document r : result)
        System.out.println(r);}
```

# Complex Queries and Aggregation with MapReduce

---

Document{[\_id=action, value=Document{{average=398.9032258064516, count=31.0}}}]}

Document{[\_id=detective, value=Document{{average=597.1724137931035, count=29.0}}}]}

Document{[\_id=drama, value=Document{{average=536.88, count=25.0}}}]}

Document{[\_id=educational, value=Document{{average=536.1739130434783, count=23.0}}}]}

Document{[\_id=fantasy, value=Document{{average=540.0434782608696, count=23.0}}}]}

Document{[\_id=humor, value=Document{{average=603.5357142857143, count=28.0}}}]}

Document{[\_id=romance, value=Document{{average=497.39285714285717, count=28.0}}}]}

Document{[\_id=thriller, value=Document{{average=513.5862068965517, count=29.0}}}]}

# SQL After All

---

- GROUP BY-style SQL queries are convertible to an equivalent map–reduce pipeline
- Many document store implementations express queries using an SQL interface
- Couchbase also allows defining foreign keys and performing join operations

```
SELECT books.title, books.genres,  
authors.name
```

```
FROM books
```

```
JOIN authors ON KEYS books.authorId
```

# SQL After All

---

- Many RDBMS vendors start implementing NoSQL by the following:
  - Focusing on horizontal scalability and distributed querying
  - Dropping schema requirements
  - Support for nested data types or allowing storing JSON directly in tables
  - Support for map–reduce operations
  - Support for special data types, such as geospatial data

# Column-Oriented Databases

---

- A **column-oriented DBMS** is a database management system that stores data tables as sections of columns of data
- Useful if:
  - Aggregates are regularly computed over large numbers of similar data items
  - Data are sparse, i.e., columns with many null values
- Can also be an RDBMS, key–value, or document store

# Column-Oriented Databases

---

- Example

Id	Genre	Title	Price	Audiobook price
1	fantasy	My first book	20	30
2	education	Beginners guide	10	null
3	education	SQL strikes back	40	null
4	fantasy	The rise of SQL	10	null

- Row-based databases are not efficient at performing operations that apply to the entire dataset
  - Need indexes which add overhead

# Column-Oriented Databases

---

- In a column-oriented database, all values of a column are placed together on disk

Genre: fantasy:1,4 education:2,3

Title: My first book:1      Beginners guide:2   SQL strikes back:3   The rise of SQL:4

Price: 20:1      10:2,4      40:3

Audiobook price: 30:1

- A column matches the structure of a normal index in a row-based system
- Operations such as find all records with price equal to 10 can now be executed directly
- Null values do not take up storage space anymore



# Column-Oriented Databases


---

- Disadvantages
  - Retrieving all attributes pertaining to a single entity becomes less efficient
  - Join operations will be slowed down
- Examples
  - Google BigTable, Cassandra, HBase, and Parquet

# Graph-Based Databases

---

- **Graph databases** apply graph theory to the storage of information of records
- Graphs consist of **nodes** and **edges**



# Graph-Based Databases

---


- One-to-one, one-to-many, and many-to-many structures can easily be modeled in a graph
- Consider the N–M relationship between books and authors
- RDBMS needs three tables: Book, Author and Books\_Authors
- SQL query to return all book titles for books written by a particular author would look like this:

```
SELECT title
FROM books, authors, books_authors
WHERE author.id = books_authors.author_id
      AND books.id = books_authors.book_id
      AND author.name = "Bart Baesens"
```

# Graph-Based Databases

---

- In a graph database (using **Cypher query language** from Neo4j)



```
MATCH (b:Book)<-[:WRITTEN_BY]-(a:Author)
WHERE a.name = "Bart Baesens"
RETURN b.title
```

# Graph-Based Databases

---

- A graph database is a hyper-relational database, in which JOIN tables are replaced by more interesting and semantically meaningful relationships that can be navigated and/or queried using graph traversal based on graph pattern matching.

# Graph-Based Databases

---

- Cypher Overview (Neo4j)
- Exploring a social graph

# Cypher Overview

---

- Cypher is a declarative, text-based query language, containing many similar operations as SQL
- Contains a special **MATCH** clause to match those patterns using symbols that look like graph symbols as drawn on a whiteboard
- Nodes are represented by parentheses, representing a circle: ( )
- Nodes can be labeled in case they need to be referred to elsewhere, and be further filtered by their type, using a colon: (b:Book)
- Edges are drawn using either - - or - - >, representing a unidirectional line or an arrow representing a directional relationship, respectively

# Cypher Overview

---

- Relationships can be filtered by putting square brackets in the middle:  
`(b:Book) <- [:WRITTEN_BY] - (a:Author)`



# Cypher Overview

---

```
MATCH (b:Book)
```

```
RETURN b;
```

```
MATCH (b:Book)
```

```
RETURN b
```

```
ORDER BY b.price DESC
```

```
LIMIT 20;
```

```
MATCH (b:Book)
```

```
WHERE b.title = "Beginning Neo4j"
```

```
RETURN b;
```

```
MATCH (b:Book {title:"Beginning Neo4j"})
```

```
RETURN b;
```

# Cypher Overview

---

- JOIN clauses are expressed using direct relational matching

```
MATCH (c:Customer)-[p:PURCHASED]->(b:Book)<-  
[:WRITTEN_BY]-(a:Author)  
WHERE a.name = "Wilfried Lemahieu"  
      AND c.age > 30  
      AND p.type = "cash"  
RETURN DISTINCT c.name;
```

# Cypher Overview

---

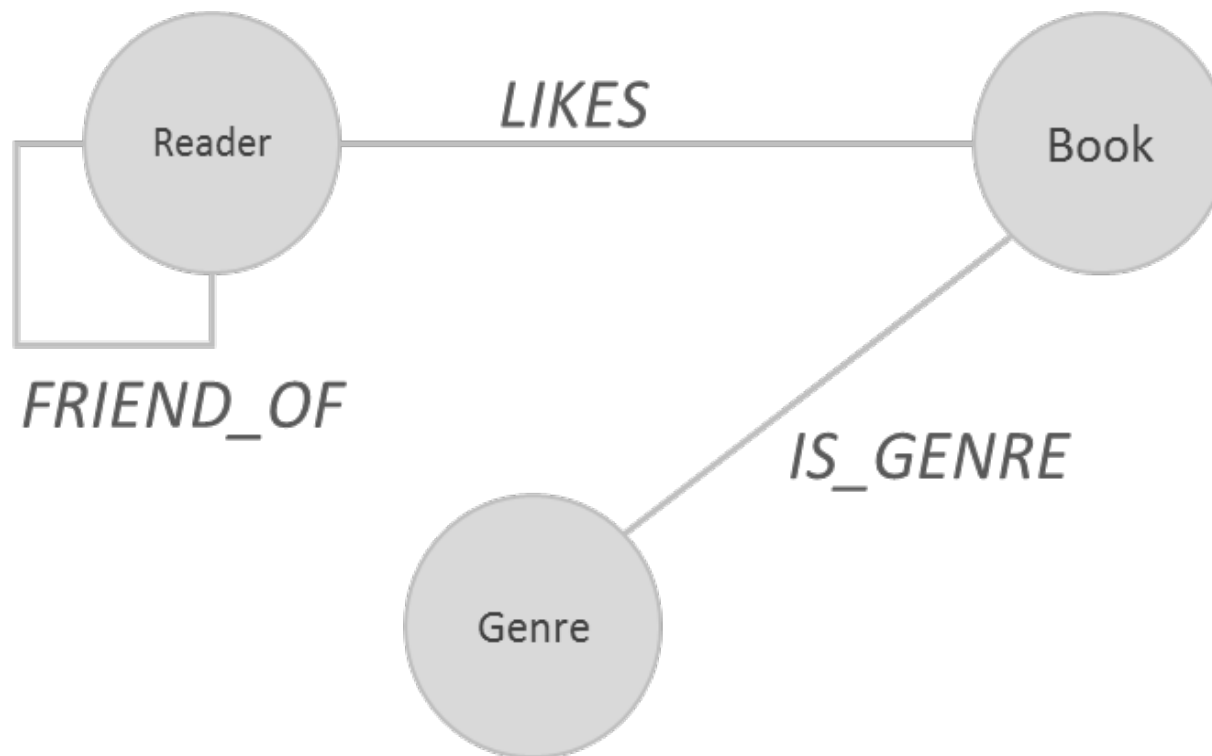
- Graph databases are great at managing tree structures
- Example:
  - A tree of book genres and books can be placed under any category level
  - A query to fetch a list of all books in the category “Programming” and all its subcategories
- Cypher can express queries over hierarchies and transitive relationships of any depth simply by appending an asterisk \* after the relationship type and providing optional min..max limits

```
MATCH (b:Book)-[:IN_GENRE]->(:Genre)  
    -[:PARENT*0..]-(:Genre {name:"Programming"})  
RETURN b.title;
```

# Exploring a Social Graph

---

- Example: a social graph for a book-reading club, modeling genres, books, and readers



# Exploring a Social Graph

---

```
CREATE (Bart:Reader {name:'Bart Baesens', age:32})
CREATE (Seppe:Reader {name:'Seppe vanden Broucke', age:30})
...

CREATE (Fantasy:Genre {name:'fantasy'})
CREATE (Education:Genre {name:'education'})
...

CREATE (b01:Book {title:'My First Book'})
CREATE (b02:Book {title:'A Thriller Unleashed'})
...

CREATE
  (b01)-[:IS_GENRE]->(Education),
  (b02)-[:IS_GENRE]->(Thriller),
...

CREATE
  (Bart)-[:FRIEND_OF]->(Seppe),
  (Bart)-[:FRIEND_OF]->(Wilfried),
...

CREATE
  (Bart)-[:LIKES]->(b01), (Bart)-[:LIKES]->(b03),
  (Bart)-[:LIKES]->(b05), (Bart)-[:LIKES]->(b06),
...
```

[illegible]

# Exploring a Social Graph

---

- Who likes romance books?

```
MATCH (r:Reader) -- (:Book) -- (:Genre  
{name: 'romance'})
```

```
RETURN r.name
```

Returns:

Elvis Presley

Mike Smith

Anne HatsAway

Robert Bertoli

...

# Exploring a Social Graph

---

- Who are Bart's friends that liked Humor books?

```
MATCH (me:Reader)--(friend:Reader)--(b:Book)--(g:Genre)
WHERE g.name = 'humor' AND me.name = 'Bart Baesens'
RETURN DISTINCT friend.name
```

- Can you recommend some humor books that Seppe's friends liked and Seppe has not liked yet?

```
MATCH (me:Reader)--(friend:Reader),
      (friend)--(b:Book),
      (b)--(genre:Genre)
WHERE NOT (me)--(b)
      AND me.name = 'Seppe vanden Broucke' AND genre.name = 'humor'
RETURN DISTINCT b.title
```



# Exploring a Social Graph

---

- Get a list of people who have liked books Bart liked, sorted by most liked books in common

```
MATCH (me:Reader)--(b:Book),  
      (me)--(friend:Reader)--(b)  
WHERE me.name = 'Bart Baesens'  
RETURN friend.name, count(*) AS common_likes  
ORDER BY common_likes DESC
```

friend.name	common_likes
Wilfried Lemahieu	3
Seppe vanden Broucke	2
Mike Smith	1

# Graph Databases

---

- Location-based services
- Recommender systems
- Social media (e.g., Twitter and FlockDB)
- Knowledge-based systems

# Other NoSQL Categories

---

- XML databases
- OO databases
- Database systems to deal with time series and streaming events
- Database systems to store and query geospatial data
- Database systems such as BayesDB which let users query the probable implication of their data

# Evaluating NoSQL DBMSs

---

- Most NoSQL implementations have yet to prove their true worth in the field
- Some queries or aggregations are particularly difficult; map–reduce interfaces are harder to learn and use
- Some early adopters of NoSQL were confronted with some sour lessons
  - e.g., Twitter and HealthCare.gov

# Evaluating NoSQL DBMSs

---

- NoSQL vendors start focusing again on robustness and durability, whereas RDBMS vendors start implementing features to build schema-free, scalable data stores
- NewSQL: blend the scalable performance and flexibility of NoSQL systems with the robustness guarantees of a traditional RDBMS

# Evaluating NoSQL DBMSs

	<b>RDBMSs</b>	<b>NoSQL Databases</b>	<b>NewSQL</b>
<b>Relational</b>	Yes	No	Yes
<b>SQL</b>	Yes	No	Yes
<b>Column stores</b>	No	Yes	Yes
<b>Scalability</b>	Limited	Yes	Yes
<b>Eventually consistent</b>	Yes	Yes	Yes
<b>BASE</b>	No	Yes	No
<b>Big volumes of data</b>	No	Yes	Yes
<b>Schema-less</b>	No	Yes	No

# Conclusion

---

- The NoSQL movement
- Key–value stores
- Tuple and document stores
- Column-oriented databases
- Graph-based databases
- Other NoSQL categories