

Flume

# Introduction

- Distributed, reliable system for efficient collection, aggregation and movement of streaming data
- Common uses
  - moving log data
  - event data
    - Social media feeds
    - Message event queues
    - Network traffic data

- Flume sources  
consume events from external sources  
forward it to channels

- External sources  
any system that generates events
  - Social media feed
    - Twitter
    - Machine logs
    - Message queues

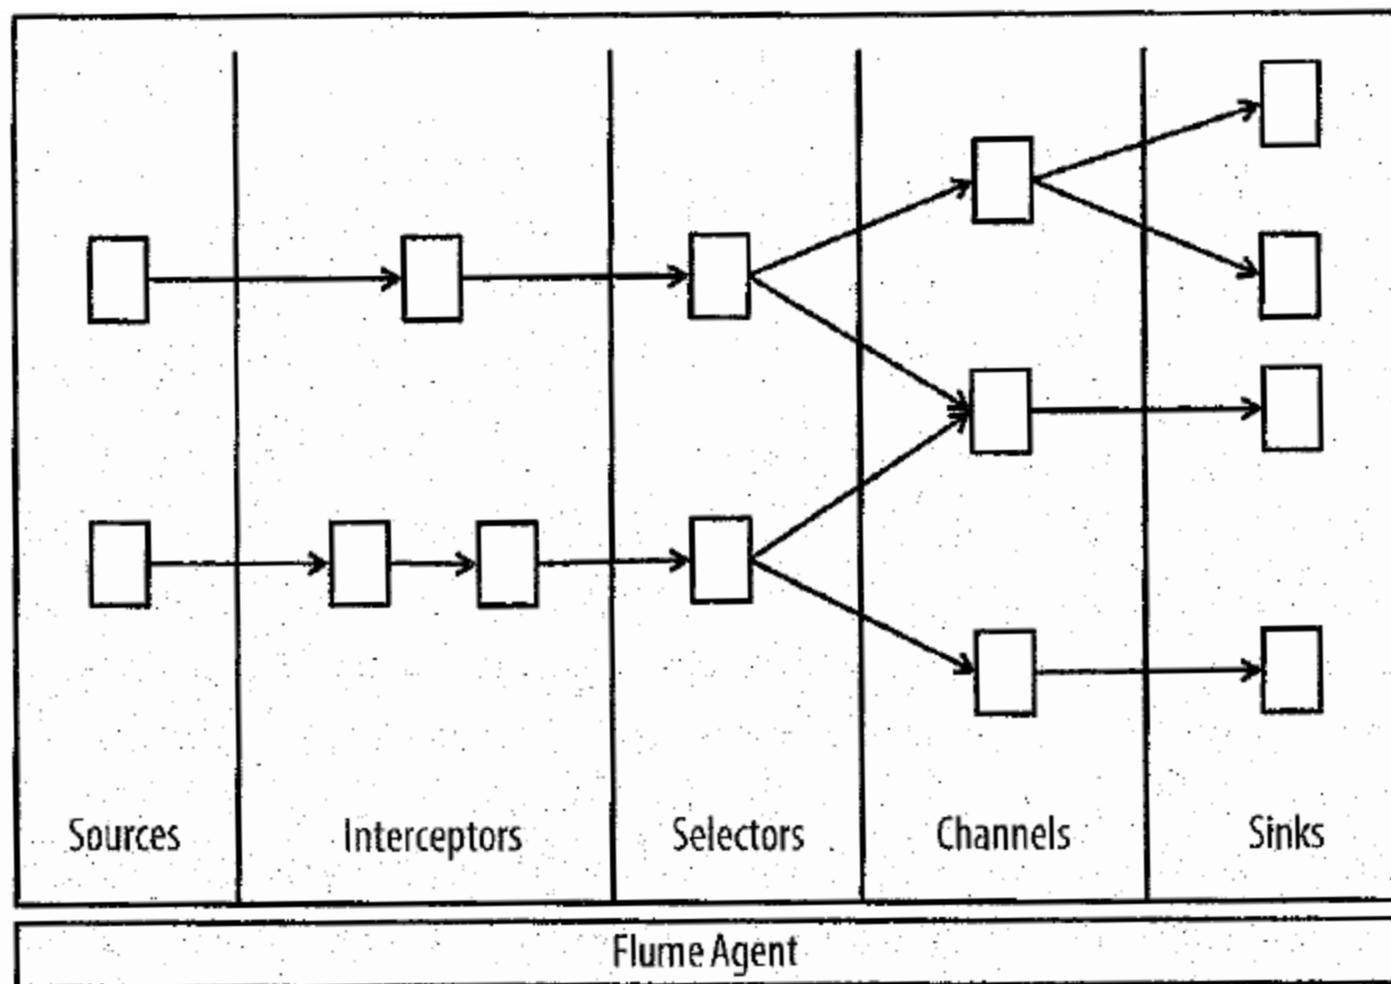
- Flume interceptors

allow events to be intercepted and modified in flight

- Transforming the event
- Enriching the event
- Anything that can be implemented in a java class

examples

- Formatting
- Partitioning
- Filtering
- Splitting
- Validating
- ...



- Selectors

provide routing for events

to send events down 0 or more paths

- To fork to multiple channels
- Send to a specific channel based on event

- Channels

store events until they are consumed by a sink

common channels

- Memory channel
- File or disk channel

## Memory channel

- Stores events in memory
- Best performance
- Least reliable as events will be lost if process or host goes down

## Disk channel

- Durable storage of events by persisting to disk

- Sinks


remove events from a channel and deliver it to a destination  
destination

- Final target for events
- Feed into further Flume processing

## Example

- HDFS sink – writes events into HDFS files






- Agent  
container

JVM process hosting of a set of Flume:

- Sources
- Sinks
- Channels
- ...



# Flume features and patterns

- Reliability
  - events stored in channel until delivered to next stage
- Recoverable
  - events persisted to disk and recovered in event of failure
- Declarative
  - no coding
    - Configuration specifies how components are wired together
- Highly customizable
  - pluggable architecture

- Fan-in

flume agent on each source system (say web servers)

- agent sends events to agents on Hadoop edge nodes

- Edge nodes

nodes on Hadoop cluster

- Features

multiple edge nodes


- Reliability – if one edge node goes down, events not lost

compress events to reduce traffic

SSL to encrypt data

- Splitting data on ingest  
split events for ingestion into multiple targets  
example
  - Send events into a primary cluster and a backup cluster

- Partitioning data on ingest  
partition data as it is ingested  
example
  - Partition events by timestamp





- Flume source consumes events delivered to it by an external source like a web server.
- External source sends events to Flume in a format that is recognized by the target Flume source.
  - Example - Avro Flume source can receive Avro events from Avro clients or
  - other Flume agents in the flow that send events from an Avro sink.
- When Flume source receives an event, it stores it into one or more channels.
- Channel - passive store that keeps the event until it's consumed by a Flume sink.
  - Example - file channel – backed by the local filesystem.

- Sink removes event from the channel
  - puts it into an external repository like HDFS (via Flume HDFS sink) or
  - forwards it to the Flume source of the next Flume agent (next hop) in the flow.
- Source and sink within agent run asynchronously with events staged in the channel.

- Splitting events for streaming analytics

sending to a streaming analytics engine such as storm or spark streaming

- Real-time counts
- Windowing
- Summaries

spark streaming implement interface for Flume Avro source

- Point a Flume Avro sink to Spark Streaming's Flume Stream

# File Formats

- Text files
  - not optimal for HDFS
  - better to save to SequenceFiles (default for HDFS sink)
  - or save to Avro
- Columnar formats
  - RCFile, ORC, Parquet

# Best Practices – Flume sources

- Batch size

Example

Client sends batch of events to Avro source

Client must wait until those events are in the channel and Avro source has responded back to client with an acknowledgment of success

- Latency can be an issue
- Select appropriate batch size

- Threads

Avro source

Can have many connections to it at the same time  
to make Avro source multi-threaded

- Add more clients or client threads

JVM source

pull source

to get more threads

- Configure more sources in Flume agent

# Best Practices – Flume sinks

- Number of sinks

sink can fetch data from a single channel

many sinks can fetch data from that same channel

sink runs in a single thread

- Example

HDFS gives 30Mbps to a single sink

therefore 30Mbps throughput only

more sinks consuming from the same channel will resolve this bottleneck

- Batch sizes

if sink is getting small batches, lot of time lost to executing system calls such as fsync



# Best Practices – Channels

- Memory channels

limit number of memory channels on a single node

More memory channels on a single node – less memory available to each of these channels

memory channel can be fed by multiple sources and be fetched from by multiple sinks

- File channels

multiple file channels write to multiple disks

# Flume Configuration File

- Name the components of the current agent.
- Describe/Configure the source.
- Describe/Configure the sink.
- Describe/Configure the channel.
- Bind the source and the sink to the channel.
- Can have multiple agents in Flume.
  - Differentiate each agent by using a unique name.
  - Using this name, we have to configure each agent.

- Name/list components such as sources, sinks, and the channels of the agent

```
agent_name.sources = source_name
```

```
agent_name.sinks = sink_name
```

```
agent_name.channels = channel_name
```

```
TwitterAgent.sources = Twitter  
TwitterAgent.channels = MemChannel  
TwitterAgent.sinks = HDFS
```

Transferring Twitter data using Twitter source through a memory channel to an HDFS sink, and the agent name id TwitterAgent

# Source

- Each source will have a separate list of properties.
- Property named “type” is common to every source, and it is used to specify the type of the source we are using.
- Provide the values of all the **required** properties of a particular source to configure it

```
agent_name.sources. source_name.type = value
```

```
agent_name.sources. source_name.property2 = value
```

```
agent_name.sources. source_name.property3 = value
```

```
TwitterAgent.sources.Twitter.type = Twitter (type name)
TwitterAgent.sources.Twitter.consumerKey =
TwitterAgent.sources.Twitter.consumerSecret =
TwitterAgent.sources.Twitter.accessToken =
TwitterAgent.sources.Twitter.accessTokenSecret =
```



# Sink

- Each sink will have a separate list of properties.
- Property named “type” - specifies the type of the sink
- Provide values to all the **required** properties of a particular sink

```
agent_name.sinks. sink_name.type = value
```

```
agent_name.sinks. sink_name.property2 = value
```

```
agent_name.sinks. sink_name.property3 = value
```

# HDFS sink

```
TwitterAgent.sinks.HDFS.type = hdfs (type name)
```

```
TwitterAgent.sinks.HDFS.hdfs.path = HDFS directory's  
Path to store the data
```

# Channels

- Channels to transfer data between sources and sinks.
- Need to describe the channel used in the agent.
- To describe each channel - set the required properties

```
agent_name.channels.channel_name.type = value  
agent_name.channels.channel_name. property2 = value  
agent_name.channels.channel_name. property3 = value
```

# Memory channel

```
TwitterAgent.channels.MemChannel.type = memory (type name)
```

# Binding source and sink to channel

- Since the channels connect the sources and sinks, it is required to bind both of them to the channel

```
agent_name.sources.source_name.channels = channel_name  
agent_name.sinks.sink_name.channels = channel_name
```

# twitter source, memory channel, HDFS sink.

```
TwitterAgent.sources.Twitter.channels = MemChannel
```

```
TwitterAgent.sinks.HDFS.channels = MemChannel
```



A sample configuration file with file extension .conf is shown below. It shows all the keys and keywords to be used to collect the twitter data.

```
TwitterAgent.sources = Twitter
```

```
TwitterAgent.channels = MemChannel
```

```
TwitterAgent.sinks = HDFS
```

```
TwitterAgent.sources.Twitter.type = com.cloudera.flume.source.TwitterSource
```

```
TwitterAgent.sources.Twitter.channels = MemChannel
```

```
TwitterAgent.sources.Twitter.consumerKey =
```

```
TwitterAgent.sources.Twitter.consumerSecret =
```

```
TwitterAgent.sources.Twitter.accessToken =
```

```
TwitterAgent.sources.Twitter.accessTokenSecret =
```

```
TwitterAgent.sources.Twitter.keywords = Keywords to be specified here
```

```
TwitterAgent.sinks.HDFS.channel = MemChannel
```

```
TwitterAgent.sinks.HDFS.type = hdfs
```

```
TwitterAgent.sinks.HDFS.hdfs.path = Configuration File Path to store the data. # hdfs://hadoop1:9000/rramine/Food_data/
```

```
TwitterAgent.sinks.HDFS.hdfs.fileType = DataStream
```

```
TwitterAgent.sinks.HDFS.hdfs.writeFormat = Text
```

```
TwitterAgent.sinks.HDFS.hdfs.batchSize = 100
```

```
TwitterAgent.sinks.HDFS.hdfs.rollSize = 0
```

```
TwitterAgent.sinks.HDFS.hdfs.rollCount = 0
```

```
TwitterAgent.channels.MemChannel.type = memory
```

```
TwitterAgent.channels.MemChannel.capacity = 10000
```

```
TwitterAgent.channels.MemChannel.transactionCapacity = 10000
```